

Incorporating Rich Social Interactions Into MDPs

Ravi Tejwani^{1*}, Yen-Ling Kuo^{1*}, Tianmin Shu², Bennett Stankovits¹
Dan Gutfreund³, Joshua B. Tenenbaum², Boris Katz¹, Andrei Barbu¹

Abstract—Much of what we do as humans is engage socially with other agents, a skill that robots must also eventually possess. We demonstrate that a rich theory of social interactions originating from microsociology and economics can be formalized by extending a nested MDP where agents reason about arbitrary functions of each other’s hidden rewards. This extended Social MDP allows us to encode the five basic interactions that underlie microsociology: cooperation, conflict, coercion, competition, and exchange. The result is a robotic agent capable of executing social interactions zero-shot in new environments; like humans it can engage socially in novel ways even without a single example of that social interaction. Moreover, the judgments of these Social MDPs align closely with those of humans when considering which social interaction is taking place in an environment. This method both sheds light on the nature of social interactions, by providing concrete mathematical definitions, and brings rich social interactions into a mathematical framework that has proven to be natural for robotics, MDPs.

I. INTRODUCTION

Endowing robots with the ability to understand and engage in social interactions is key to having them integrate into our daily lives. Yet, very little is known about social interactions, what they are, how to measure them, and how to computationally implement them. Research in psychology has attempted to define different types of social interactions and to test the abilities of humans to engage in them [1–4]. While no single framework for what social interactions are has emerged out of this work, one coherent proposal comes out of the field of microsociology; the study of fine-grained face-to-face interactions. Microsociology puts forward that underlying social interactions are five fundamental abilities that are recombined to give rise to the repertoire of behaviors we see in humans and other animals: cooperation, conflict, competition, coercion, and exchange. We provide the first computational mechanism for implementing these five interactions. Moreover, we put forward the first clear mathematical definition of what these five types of interactions are. The result is a principled computational mechanism for social interactions that allows robots to execute such interactions zero-shot in novel environments.

To implement these interactions, we build on top of Tejwani et al. [5], which recently introduced Social MDPs — an attempt at extending the MDP framework to social interactions. That work builds on an analogy, that the understanding of

partial observability took a great leap forward with the advent of POMDPs, and that perhaps doing the same for social interactions by building Social MDPs will similarly pay off. As executing MDPs is fairly efficient, this also ensures that the resulting models are tractable. Unfortunately, as originally formulated, Social MDPs are fundamentally limited to two of the five interactions: cooperation (helping) and conflict (hindering). Our new model is an extended Social MDP capable of handling these five fundamental social interactions. One feature of this model is that it degenerates to the original Social MDP model for the two interactions that the two share, and furthermore, it degenerates to a well-known social model in game theory for those two interactions [6].

We make four contributions:

- 1) A novel Social MDP that allows robots to execute five fundamental social interactions from microsociology.
- 2) The first formalization of what those interactions are.
- 3) A computational implementation that enables robots to engage in social interactions zero-shot, without any social-interaction-specific training.
- 4) Extensive human experiments that validate the model and demonstrate its ability to capture human judgments about social interactions.

II. RELATED WORK

In multi-agent settings, an agent learns to reason about the goals, preferences and beliefs of the other agents so that it can effectively interact with them [7, 8]. Several types of models have been explored in this space: theory-of-mind-based models for goal inference [9–13], Bayesian inverse planning [14, 15], and learning the reward functions of other agents [16]. Xie et al. [17] provide a method for learning a low-dimensional representation of the strategy of another agent. This representation enables agents to avoid or work with one another. These methods have all been limited to social interactions exhibiting cooperation or conflict (helping or hindering). Additionally, the extended Social MDPs we present are zero-shot, while most of these prior approaches require social-interaction-specific training data. We go well beyond such models supporting a far richer theory of social interactions.

Social actions such as walking, waving, hugging, and hand-shaking in videos of group activities have also been explored [18–20]. These methods broadly involve two phases [21]: a social perception phase and a coordination phase where agents interact. In contrast, Social MDPs are agnostic to specific social actions. We are not detecting if two agents are hugging and then inferring that they are friendly

* Equal contribution

¹ CSAIL & CBMM, MIT

{tejwanir, ylkuo, bstankov, boris, abarbu}@mit.edu

² BCS & CBMM, MIT {tshu, jbt}@mit.edu

³ MIT-IBM Watson AI Lab dgutfre@us.ibm.com

because we have seen that hugging often results in positive actions. We are specifying a reward for what cooperative actions are; given a novel action, that has never been seen before, in a novel context, Social MDPs determine what if any social interaction is at play.

In game theory, several approaches [6, 22] explore altruistic and spiteful actions (the term of art for cooperation and conflict) by means of linear combinations of pay-offs. Work of Levine [6] is mathematically equivalent to the level 1 Social MDPs defined in Teiwani et al. [5], while level 2 Social MDPs seem to have no counterpart in the current economics literature. This is the equivalent to saying that in a game you can help another player, but you cannot help another player whose goal is to then help a third player. Our long-term hope is to connect the extended Social MDPs presented here back to the economics literature and expand both the depth of inference considered there and the range of social interactions.

Interactive POMDPs [23–25] (I-POMDPs) are the original blueprint for Social MDPs. They extend POMDPs to allow agents to reason about other agent’s beliefs. They do not allow agents to reason about other agent’s reward functions, making it impossible for them to represent social interactions.

Human social perception has been studied extensively, but with limited theoretical insights. A rich history of work with animations of simple geometrical objects [26, 27] shows that humans readily assign social goals and intentions to any agent. More recently, agents’ behaviors have been simulated with modern physics engines in fully observable [28, 29] and partially observable environments [30, 21] to create various social benchmarks, although these benchmarks are almost exclusively limited to helping and hindering. We hope that by providing formal definitions for what social interactions are, we can eventually bring about a much closer connection between the cognitive science and robotics of social interactions.

III. MODEL

We adopt the setting of Teiwani et al. [5] who initially develop Social MDPs by giving agents a combination of physical goals and social goals. A physical goal is precisely what an MDP can represent, a reward that is a function of the state of the world — Social MDPs degenerate into MDPs when the social goal is nil. A social goal extends the reward function of every agent to allow it to contain an estimate of the reward function of another agent. In the original Social MDP formulation, every agent estimated the rewards of every other agent. Then, the reward of each agent was a linear combination of its own physical goal and its estimate of the reward function of other agents (the social goal).

The coefficient of this linear combination of physical and social goals determined how willing the agent was to engage socially. If it was zero, no weight was put on the social goal, so the agent behaved as an MDP. A large positive value gives an agent a high reward when another agent maximizes its own reward — the result is an agent that helps the other agent succeed. A large negative value does

the opposite, gives an agent a low reward when another agent maximizes its own reward — the result is an agent that stops the other agent from succeeding. Teiwani et al. [5] demonstrated that this formulation gives rise to robots that appear to humans to behave socially, and that this coefficient determined the strength of that interaction and gave rise to rich social behaviors.

Unfortunately, this framework lacks any degrees of freedom that would allow it to represent any other social interactions aside from helping and hindering. It has a single coefficient that determines the interaction between agents and the polarity of that coefficient determines if the interaction is helpful or unhelpful.

Next, we survey a precise mathematical definition of the five social interaction types considered in this framework of recursively estimating other agent’s social reward functions and then demonstrate how to extend Social MDPs to enable robots to execute these interactions.

A. The five interaction types

To define the pairwise interactions between agent i and j let the reward of an agent i at level l , R_i^l , be a combination of that agent’s physical goal, g_i and that agent’s social goal, ξ_{ij}^l ; $R_i^l = r(g_i) + \xi_{ij}^l$. Let ξ_{ij}^l be the social goal of agent i toward agent j at level l . At $l = 0$, no agents are social, the Social MDP degenerates into a traditional MDP, i.e., $\xi_{ij}^0 = 0$. At $l = 1$, each agent has a social goal, but it does not consider other agents to have social goals, only physical goals. At $l = 2$, agents assume that other agents also have social goals, etc. This nesting allows Social MDPs to start with an arbitrary level, then eventually bottom out in an MDP. Next, we will show how the reward function of agent i is formulated when that agent engages in each of the five social interaction types.

We define all interactions from the perspective of agent i , i.e., if agent i wants to be social with respect to agent j in each of these five ways, what is i ’s reward function? Reward functions are specified in terms of estimated rewards, as i does not know what j wants, it must infer j ’s reward function; we denote all estimated quantities with a tilde, as in $\tilde{R}_j^{l-1,i}$ for the estimate of j ’s reward function made by i at level $l-1$ and $\tilde{\xi}_{ij}^{l,i}$ for the estimate of the social interaction between i and j at level l again made by i . In each case, the additional superscript i denotes the agent performing the estimation — agents i and k may, depending on their observations or biases, not estimate the same reward function for agent k .

An overview of the recursive inferences being made between two agents is shown in Fig. 1. In Table I, we show the social component of each of the reward functions of agent i as a function of its own social goal and the estimated social goal of j when $l > 1$ (when other agents are assumed to have social goals) and when $l = 0$ (when other agents have only physical goals). For each agent, we estimate what kind of social interaction it is engaging in. Then, the reward function of the agent is updated according to Fig. 1. For example, if level 2 agent i intends to cooperate with j , the original reward function template for i , $R_i^l = r(g_i) + \xi_{ij}^l$, is updated




Fig. 1: Reasoning with an extended Social MDP rendered as Bayesian network at different levels. s are observations, ψ is the policy, g are physical goals, and ξ are social goals. The yellow robot, agent 1, is reasoning at level two about the red robot's (agent 2) social actions. This demonstrates the utility of nesting Social MDPs, shallower social MDPs are more limited in their reasoning about the social abilities of other agents.

Algorithm 1 The algorithm to compute social policy ψ_i^l for agent i at level l and time t . We use the estimated social policy $\tilde{\psi}_j^{l-1,i}$ at previous time step to update the estimated physical and social reward as described in Section III-C. At $t = 0$, we assume $P(\tilde{g}_j^{l,i,t})$ and $P(\tilde{\xi}_{ji}^{l,i,t})$ are sampled from uniform distributions. This algorithm is called at all recursion steps $\tilde{\psi}_j^{l-1,i}$ to estimate social policy for the other agent j . The estimated reward and policy are used to compute the Q values for selecting the actions.

Require: $l, s^t, a_i^t, a_j^t, \xi_{ij}, g_i$

if $l = 0$ **then**
 solve MDP for agent i

else
 for all $\tilde{\xi}_{ji}^{l,i,t}, \tilde{g}_j^{l,i,t}$ **do**
 compute
 $P(\tilde{\xi}_{ji}^{l,i,t} | s^{t-1}, a_i^{t-1}, a_j^{t-1})$
 $P(\tilde{g}_j^{l,i,t} | s^{1:t-1})$
 $\tilde{\psi}_j^{l-1,i}(s^t, a_j^t, a_i^t, \tilde{\xi}_{ji}^{l,i,t}, \tilde{g}_j^{l,i,t})$
 end for
 compute $R_i^l(s^t, a_i^t, a_j^t, \xi_{ij}, g_i)$
 compute $Q_i^l(s^t, a_i^t, a_j^t, \xi_{ij}, g_i)$
 $\pi_i^l \leftarrow \operatorname{argmax}_{a_i \in \mathcal{A}_i} Q_i^l$

end if

Type of ξ_{ij}^l	Type of $\tilde{\xi}_{ji}^{l,i}$	Substitute ξ_{ij}^l at $l > 1$	Substitute ξ_{ij}^l at $l = 1$
cooperation	any	$\tilde{R}_j^{l-1,i}$	$r(\tilde{g}_j^{l,i})$
conflict	any	$-\tilde{R}_j^{l-1,i}$	$-r(\tilde{g}_j^{l,i})$
competition	cooperation	$\tilde{R}_j^{l-1,i}$	$-r(\tilde{g}_j^{l,i})$
	conflict	$-\tilde{R}_j^{l-1,i}$	
	competition	$-\tilde{R}_j^{l-1,i}$	
	coercion	$\tilde{R}_j^{l-1,i}$ if $g_i = \tilde{g}_j^{l,i}$ $-\tilde{R}_j^{l-1,i}$ if $g_i \neq \tilde{g}_j^{l,i}$	
coercion	exchange	$\tilde{R}_j^{l-1,i}$	$r(g_i)$
	cooperation	$\tilde{R}_j^{l-1,i} + r(g_i)$	
	conflict	$-\tilde{R}_j^{l-1,i} + r(g_i)$	
	competition	$\tilde{R}_j^{l-1,i}$ if $g_i = \tilde{g}_j^{l,i}$ $-\tilde{R}_j^{l-1,i} + r(g_i)$ if $g_i \neq \tilde{g}_j^{l,i}$	
	coercion	$-\tilde{R}_j^{l-1,i} + r(g_i)$	
exchange	exchange	$\tilde{R}_j^{l-1,i} + r(g_i)$	N/A
	others	$n \cdot r(\tilde{g}_j^{l,i}) + \tilde{R}_j^{l-1,i}$ N/A	

TABLE I: Computing the reward function of a social agent i , engaging with agent j in one of five social interactions at level l . The reward function for i is of the form $R_i^l = r(g_i) + \xi_{ij}^l$, where $r(g_i)$ is the reward for the physical goal of agent i (the goal and reward in a plain MDP) and ξ_{ij}^l is the social interaction of i toward j . Based on the type of interaction that i wants to have with j and its estimate of what the reciprocal interaction is (i 's estimate of how j wants to interact with i , $\tilde{\xi}_{ji}^{l,i}$) we show what term to substitute ξ_{ij}^l with achieve that interaction. For example, in cooperation, we substitute i 's estimate of what j wants into i 's reward function; this gives i a reward whenever j accomplishes their goals. N/A are impossible interactions. See the text for an explanation of the other substitutions.

as $R_i^l = r(g_i) + \tilde{R}_j^{l-1,i}$ — agent i 's reward now includes agent j 's reward so i will help j .

Cooperation and *Conflict*, the two types of goals supported by the original Social MDPs are peculiar: the reward function does not depend on what kind of social interaction the other agent wants to engage in. This is why they could be implemented in the original Social MDP formulation. The other three types of social interaction are more complex, they depend on the other agent's intentions. For example, competing with someone who intends to cooperate with you is very different from competing with someone who intends to stop you from taking actions.

See Table I for the complete mapping from interactions to rewards. *Competition* and *Coercion* follow a pattern that is similar to cooperation and conflict with several special cases since they are situation-specific. In other words, what it means to compete depends on what the other agent wants to do, if they are pleasant you incorporate their reward differently into yours than if they are combative.

Exchange is the outlier. Exchanging implies that an agent performs an action that is useful for another and vice versa. If one of the two agents isn't willing to exchange, it's impossible for either to do so on their own hence the N/A in Table I. When directing an agent to exchange with another, we update that agent's reward function to include a copy of the other agent's goal with a positive, but small, scalar n . This makes the agent willing to help, weakly, while maintaining its own goals. If both agents do this symmetrically, the result is a pair which is willing to perform mutually-beneficial actions as long as they can accomplish their other goals more effectively.

In the results section we demonstrate that the method outlined in Table I of mapping different social interactions to reward functions gives rise to behavior which is interpreted as the correct social interaction by humans. Next we describe how to build and perform inference with an extended Social MDP, one which can handle this richer set of social interactions and which estimates the type of interaction rather than merely supporting cooperation and conflict.

B. Extended Social MDPs

A Social MDP for agent i with respect to all agents J consists of an arity (here we formulate the pairwise case) and a maximum level, l , and is defined as:

$$M_i^l = \langle \mathcal{S}, \mathcal{A}, T, \xi_{ij}, g_i, R_i^l, \gamma \rangle \quad (1)$$

where \mathcal{S} is a set of states in the environment where $s \in \mathcal{S}$; $\mathcal{A} = \mathcal{A}_i \times \mathcal{A}_j$ is the set of joint moves of agent i and agent j ; T is the probability distribution of going from state $s \in \mathcal{S}$ to next state $s' \in \mathcal{S}$ given actions of both agents: $T(s' | s, a_i, a_j)$; ξ_{ij} is agent i 's intended social interaction with agent j . It computes agent i 's social reward when interacting with agent j ; g_i is agent i 's physical goal; R_i^l is the l -th level reward function for agent i based on its estimate of other agents' rewards; and γ is a discount factor, $\gamma \in (0, 1)$.

Each agent has its own physical goal, e.g., going to a landmark, as well a social goal, e.g., helping or hindering

other agents. What enables Social MDPs to go beyond regular MDPs is the recursive nature of the reward function which can be written in terms of the estimated rewards of other agents. The immediate reward of an agent i at each time step is computed as follows:

$$R_i^l(s, a_i, a_j, \xi_i, g_i) = r(g_i) + \xi_i(g_i, \tilde{g}_j, \tilde{\xi}_j) - c(a_i) \quad (2)$$

C. Planning with Social MDPs

Analogous to MDPs, the Q function of Social MDPs is the sum of immediate reward and the expected value in the future.

$$Q_i^l(s, a_J, \xi_{ij}, g_i) = R(s, a_i, \xi_{ij}, g_i) + \gamma \sum_{s' \in \mathcal{S}} T(s, a_J, s') V_i^l(s', \xi_{ij}, g_i) \quad (3)$$

Since agent i is interacting with other agents $j \in J$, it needs to estimate what actions other agents are likely to take in order to compute its state-action value. Social MDPs take the expectation over the estimated goals and actions of agent j to compute $V_i^l(s', \xi_{ij}, g_i)$:

$$\begin{aligned} V_i^l(s', \xi_{ij}, g_i) &= \max_{a'_i \in \mathcal{A}_i} \left\{ E_{\tilde{g}_j^{l,i}, \tilde{\xi}_{ji}^{l,i}, a'_j} [Q_i^l(s', a'_J, \xi_{ij}, g_i)] \right\} \\ &= \max_{a'_i \in \mathcal{A}_i} \left\{ \sum_{j \in J, j \neq i} \sum_{a'_j \in \mathcal{A}_j} \sum_{\tilde{g}_j^{l,i}} \int_{\tilde{\xi}_{ji}^{l,i}} \underbrace{P(\tilde{g}_j^{l,i} | s^{1:t})}_{\text{estimate physical goal (Eq. 6)}} \right. \\ &\quad \left. \underbrace{P(\tilde{\xi}_{ji}^{l,i} | s, a_J)}_{\text{estimate social goal (Eq. 5)}} \underbrace{\tilde{\psi}_j^{l-1,i}(s', a'_J, \tilde{\xi}_{ji}^{l,i}, \tilde{g}_j^{l,i})}_{\text{estimate social policy (Eq. 7)}} Q_i^l(\cdot) d\tilde{\xi}_{ji}^{l,i} \right\} \end{aligned} \quad (4)$$

When solving agent i 's MDP at level l , the estimated social and physical goals are further used to update the other agent j 's social policy to the actions agent j may take. We denote the estimated social policy for agent j at reasoning level $l-1$ as $\tilde{\psi}_j^{l-1,i} : \mathcal{S} \times \mathcal{A}_J \times \tilde{\xi}_{ji}^{l,i} \times \tilde{g}_j^{l,i} \rightarrow [0, 1]$. Algorithm 1 summarizes the steps to compute the state-action values and select optimal actions for any level l at time step t . We first update the probability of the estimated goals of other agents using the observed state and the estimated policy from the previous time step. The updated probability of goals are used to update the policy of other agents and compute the reward and Q function of the target agent.

An agent's estimate of another agent's social and physical goals at time step t and level l can be updated based on the actions performed by the agents. At $t=0$, we use uniform distributions for social and physical goals. The social goal, estimated at time step t , is updated after actions taken by all agents at the previous time step. This update is similar to the belief update in the POMDP framework but based on the estimated social policy of the other agent j :

$$\begin{aligned} P(\tilde{\xi}_{ji}^{l,i,t} | s^{t-1}, a_J^{t-1}) &\propto P(\tilde{\xi}_{ji}^{l,i,t-1} | s^{t-2}, a_J^{t-2}) \\ \sum_{\tilde{g}_j^{l,i,t-1}} P(a_j^{t-1} | s^{t-1}, \tilde{\xi}_{ji}^{l,i,t-1}, \tilde{g}_j^{l,i,t-1}) &\times T(s^{t-1}, a_J^{t-1}, s^t) \end{aligned} \quad (5)$$

The physical goal g_j of agent j is estimated by i as follows, similar to [28] but marginalized over the estimated social goal as the agent is estimating the social goal at the same time.

$$P(\tilde{g}_j^{l,i,t} | s^{1:t-1}) \propto \int_{\tilde{\xi}_{ji}^{l,i,t}} P(s^{1:t-1} | \tilde{g}_j^{l,i,t}, \tilde{\xi}_{ji}^{l,i,t}) P(\tilde{g}_j^{l,i,t}) P(\tilde{\xi}_{ji}^{l,i,t}) d\tilde{\xi}_{ji}^{l,i,t} \quad (6)$$

The l -level social policy $\tilde{\psi}_j^{l,i}$ of the agent j is predicted by i using the Q-function at level $l-1$:

$$\tilde{\psi}_j^{l-1,i}(s, a_J, \tilde{\xi}_{ji}^{l,i}, \tilde{g}_j^{l,i}) = \text{Softmax}(Q_j^{l-1}(s, a_J, \tilde{\xi}_{ji}^{l,i}, \tilde{g}_j^{l,i})) \quad (7)$$

This is a softmax policy where we use a temperature parameter τ to control how much the agent j follows greedy actions. As shown in eq. 4, in order to use agent j 's Q function at level $l-1$, it requires to compute agent i 's Q function at level $l-2$, and so on. Recursively solving Social MDPs eventually bottoms out in level 0 where one solves an MDP.

IV. RESULTS

We apply our extended Social MDP framework to a multi-agent gridworld inspired by previous studies on social perception [15, 9, 31, 30]. This 10×10 world consists of two agents (a yellow robot and red robot), two physical landmarks (a construction site and a tree) and three objects (an axe, wooden log, and a water bucket). Objects can be pushed to either destination. Physical goals consist of moving the desired objects to one of the landmarks. Agents can have no social goal or one of the five social goals: cooperation, conflict, competition, coercion, or exchange. At any point in time, agents can push an object forward, move in one of the four cardinal directions or choose to take no actions.

Each agent's reward for reaching its physical goal is based on that agent's geodesic distance from the goal after taking an action [15]. This physical reward function is parameterized by ρ and δ that determines the scale and shape of the physical reward: $r(s, a, g_i) = \max(\rho(1 - \text{distance}(s, a, g_i)/\delta), 0)$. We set the cost, c , of an action a , to 1 for grid moves and to 0.1 for staying in place while ρ and δ were set to 1.25 and 5, respectively. The discount factor, γ , was set to 0.99.

We systematically enumerate the 72 unique scenarios by assigning agents either the same or different physical goal, and one of five social goals or no social goal ($2 * 6 * 6 = 72$ scenarios). Even this simple world gives rise to many rich social interactions with subtle but meaningful differences in the behaviors of agents depending on their understanding of social interactions. For instance, take the scenario where both agents have the social goal of exchanging something. An exchange is practical only in some cases: the agents must be arranged in such a way that it would be helpful for the red agent to aid the yellow agent and vice versa (for example, the red agent is closer to the axe which the yellow agent wants and the yellow agent is closer to the log which the red agent wants). Agents must recognize this. Then they must recognize if the other agent is willing to exchange by attempting the exchange. Then, they must follow through with the exchange,

bring the other object to a meeting point, swap objects, and then go on to complete their own physical goals. If one of the agents is uncooperative, both agents should abort and try to perform their own physical goals, as the exchange is then impossible. This occurs without pre-specifying any symbolic notion of "exchange", without a single training example of an exchange, and without any hardcoded rules, merely by specifying the correct reward function.

All scenarios with descriptions, diagrams, videos, and detailed timestep-by-timestep results for all human experiments and models are available at this URL

<https://social-interactions-mdp.github.io>

A. Are these social interactions?

We first establish that these scenarios do indeed show social interactions and that humans are able to recognize the intended social interactions. Each of the 72 scenarios were executed with level 2 Social MDPs guiding the yellow agent giving rise to 72 videos. 12 subjects were recruited on Mechanical Turk. Each subject saw each of the 72 videos and rated it according to their confidence that the video depicts a given social interaction after each action taken by every agent. Taking the human subject's last judgment, at which point they have seen the entire interaction, we ask how well does the most confident label that humans give to the interaction match the intended interaction from the Social MDP. Humans provided the top label with 82% accuracy, showing that Social MDPs generate the intended social interactions the vast majority of the time, see Table II.

B. Do we model human judgments?

Going further, we demonstrate that not only do humans recognize the intended social interactions from Social MDPs, but Social MDPs also predict human judgments. Taking each of the 72 videos, we compare the confidence of humans that a particular social interaction is being depicted to that of the model. We find that our model has a 0.784 correlation with human judgments (0.86 for cooperation, 0.84 for conflict, 0.64 for competition, 0.64 for coercion, 0.76 for exchange), see Fig. 2 in blue. When considering estimating the physical goal, rather than the social goal, our model has a 0.76 correlation with human judgments.

C. Models

We compare Social MDPs with two alternative models, inverse planning [15] and a recent cue-based model [28]. Each model is provided with every video, frame by frame, and is required to incrementally predict which social interaction is taking place and what the physical goal is. At our website, we list every scenario with detailed results from each model at every time step. Several example scenarios are shown in Section IV. Qualitatively, one can observe that the Social MDPs track human judgments about these videos far more accurately. Quantitatively, results aggregated across the different categories are shown in Table II. While humans are able to estimate social goals and physical goals more accurately than any model, the gap between Social MDPs and

Model	Social Goal						Physical Goal
	Cooperation	Conflict	Competition	Coercion	Exchange	Overall	
Human	0.934	0.952	0.623	0.724	0.876	0.823	0.974
Extended Social MDP (Ours)	0.845	0.851	0.471	0.651	0.814	0.726	0.831
Inverse Planning	0.763	0.784	0.261	0.283	0.197	0.457	0.783
Cue Based Model	0.461	0.434	0.127	0.156	0.083	0.252	0.432

TABLE II: The accuracy of humans and each of the models at determining which social interaction is taking place in each of the 72 scenarios. Our model is significantly more accurate, particularly when it comes to recognizing social interactions.




Fig. 2: Humans and our model scored 72 scenarios according to how likely each social interaction was and how likely one of the physical goals was. The straight line is the best linear fit and the light blue band represents the 95% confidence interval. Our model agrees with humans and predicts their confidence scores for both social interactions and the physical goal.

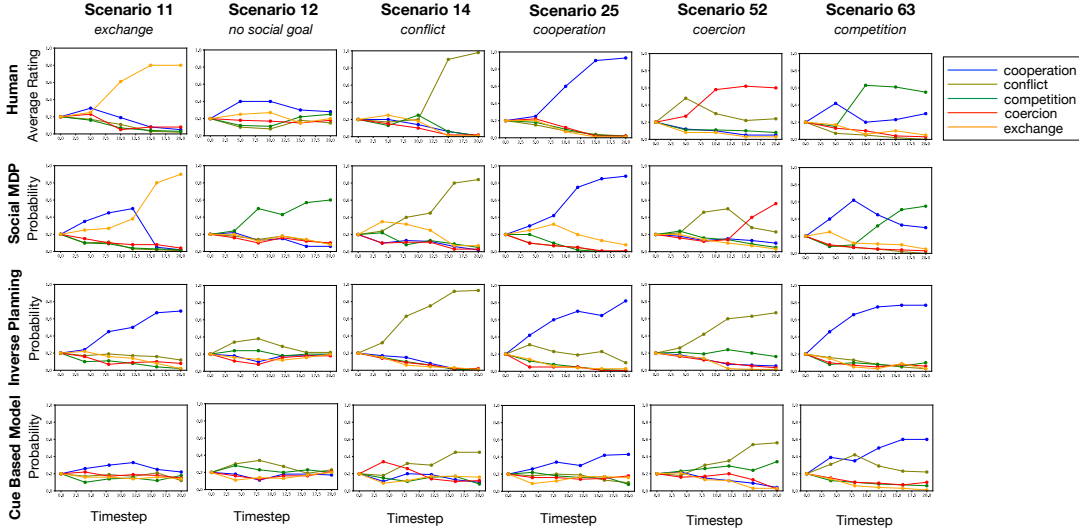


Fig. 3: Humans, Social MDPs, and the two other baseline models were asked predicted the social interaction in each scenario at every time step. We chose six representative scenarios that each show a different interaction. Above every column we provide the scenario number along with its intended social interaction (our website provides rendering of every scenario along with many other details for each). Social MDPs agree with humans timestep by timestep, with the exception of scenario 12 where humans recognize that there is no social interaction while Social MDPs are $\approx 60\%$ certain that agents are competing.

other models is stark. Not only do Social MDPs estimate the physical goal more accurately, but the social goal accuracy is dramatically higher (72.6% vs 45.7%). Our online supplement also demonstrates the different levels of nesting.

V. CONCLUSION

We demonstrated extended Social MDPs which are able to reason about the five core social interactions described in the microsociology literature. This model is capable of reasoning about social interactions without requiring interaction-specific training, just like humans are able to recognize social interactions in novel environments. The inferences the model makes both qualitatively and quantitatively match human judgments. This significantly expands the space of social

interactions that robots engage in. The runtime of the model depends on the number of levels considered as this controls how many MDPs must be solved recursively. We do not yet know which level corresponds to the inferences that humans are able to make.

We would like to scale this approach up to continuous environments and to physical robots with manipulators. In addition, we would like to recognize these social interactions from real-world videos rather than from renderings and trajectories of agents. Hopefully, sometime in the future, robots will have general-purpose social skills and the field of studying social interactions will be based on mathematical and refutable theories that span robotics, social science, cognitive science, vision, and neuroscience.

REFERENCES

- [1] E. Goffman *et al.*, *The presentation of self in everyday life*. Harmondsworth London, 1978, vol. 21.
- [2] S. N. Durlauf and Y. M. Ioannides, “Social interactions,” *Annual Review of Economics*, vol. 2, no. 1, pp. 451–478, 2010. [Online]. Available: <https://doi.org/10.1146/annurev.economics.050708.143312>
- [3] M. Weber, *The Nature of Social Action*. Cambridge University Press, 1978, p. 7–32.
- [4] E. Goffman, “On face-work: An analysis of ritual elements in social interaction,” *Psychiatry*, vol. 18, no. 3, pp. 213–231, 1955.
- [5] R. Tejwani, Y.-L. Kuo, T. Shu, B. Katz, and A. Barbu, “Social interactions as recursive mdp,” in *Conference on Robot Learning (CoRL)*, 2021.
- [6] D. K. Levine, “Modeling altruism and spitefulness in experiments,” *Review of economic dynamics*, vol. 1, no. 3, pp. 593–622, 1998.
- [7] S. V. Albrecht and P. Stone, “Autonomous agents modelling other agents: A comprehensive survey and open problems,” *Artificial Intelligence*, vol. 258, pp. 66–95, 2018.
- [8] H. He, J. Boyd-Graber, K. Kwok, and H. Daumé III, “Opponent modeling in deep reinforcement learning,” in *International conference on machine learning*. PMLR, 2016, pp. 1804–1813.
- [9] C. L. Baker, N. D. Goodman, and J. B. Tenenbaum, “Theory-based social goal inference,” in *Proceedings of the thirtieth annual conference of the cognitive science society*. Cognitive Science Society, 2008, pp. 1447–1452.
- [10] C. L. Baker and J. B. Tenenbaum, “Modeling human plan recognition using bayesian theory of mind,” *Plan, activity, and intent recognition: Theory and practice*, vol. 7, pp. 177–204, 2014.
- [11] J. Kiley Hamlin, T. Ullman, J. Tenenbaum, N. Goodman, and C. Baker, “The mentalistic basis of core social cognition: Experiments in preverbal infants and a computational model,” *Developmental science*, vol. 16, no. 2, pp. 209–226, 2013.
- [12] M. Kleiman-Weiner, M. K. Ho, J. L. Austerweil, M. L. Littman, and J. B. Tenenbaum, “Coordinate to cooperate or compete: abstract goals and joint intentions in social interaction,” in *CogSci*, 2016.
- [13] N. Rabinowitz, F. Perbet, F. Song, C. Zhang, S. A. Eslami, and M. Botvinick, “Machine theory of mind,” in *International conference on machine learning*. PMLR, 2018, pp. 4218–4227.
- [14] C. L. Baker, R. Saxe, and J. B. Tenenbaum, “Action understanding as inverse planning,” *Cognition*, vol. 113, no. 3, pp. 329–349, 2009.
- [15] T. D. Ullman, C. L. Baker, O. Macindoe, O. Evans, N. D. Goodman, and J. B. Tenenbaum, “Help or hinder: Bayesian models of social goal inference,” Massachusetts Institute of Technology Department of Brain and Cognitive Sciences, Tech. Rep., 2009.
- [16] D. Hadfield-Menell, A. Dragan, P. Abbeel, and S. Russell, “Cooperative inverse reinforcement learning,” *arXiv preprint arXiv:1606.03137*, 2016.
- [17] A. Xie, D. P. Losey, R. Tolsma, C. Finn, and D. Sadigh, “Learning latent representations to influence multi-agent interaction,” in *Conference on Robot Learning (CoRL)*, 2020.
- [18] A. Patron-Perez, M. Marszalek, I. Reid, and A. Zisserman, “Structured learning of human interactions in tv shows,” *IEEE Transactions on Pattern Analysis and Machine Intelligence*, vol. 34, no. 12, pp. 2441–2453, 2012.
- [19] M. J. Marin-Jimenez, A. Zisserman, M. Eichner, and V. Ferrari, “Detecting people looking at each other in videos,” *International Journal of Computer Vision*, vol. 106, no. 3, pp. 282–296, 2014.
- [20] M. S. Ryoo and J. K. Aggarwal, “Spatio-temporal relationship match: Video structure comparison for recognition of complex human activities,” in *2009 IEEE 12th international conference on computer vision*. IEEE, 2009, pp. 1593–1600.
- [21] X. Puig, T. Shu, S. Li, Z. Wang, J. B. Tenenbaum, S. Fidler, and A. Torralba, “Watch-and-help: A challenge for social perception and human-ai collaboration,” in *The International Conference on Learning Representations*, 2021.
- [22] Z. Feng, C. Tan, J. Zhang, and Q. Zeng, “Bargaining game with altruistic and spiteful preferences,” *Group Decision and Negotiation*, vol. 30, no. 2, pp. 277–300, 2021.
- [23] P. Doshi and P. J. Gmytrasiewicz, “Monte carlo sampling methods for approximating interactive pomdps,” *Journal of Artificial Intelligence Research*, vol. 34, pp. 297–337, 2009.
- [24] P. Doshi and D. Perez, “Generalized point based value iteration for interactive pomdps,” in *AAAI*, 2008, pp. 63–68.
- [25] B. Ng, K. Boakye, C. Meyers, and A. Wang, “Bayes-adaptive interactive pomdps,” in *Proceedings of the AAAI Conference on Artificial Intelligence*, vol. 26, no. 1, 2012.
- [26] F. Heider and M. Simmel, “An experimental study of apparent behavior,” *The American journal of psychology*, vol. 57, no. 2, pp. 243–259, 1944.
- [27] F. Abell, F. Happe, and U. Frith, “Do triangles play tricks? attribution of mental states to animated shapes in normal and abnormal development,” *Cognitive Development*, vol. 15, no. 1, pp. 1–16, 2000.
- [28] T. Shu, M. Kryven, T. D. Ullman, and J. B. Tenenbaum, “Adventures in flatland: Perceiving social interactions under physical dynamics,” in *42d proceedings of the annual meeting of the cognitive science society*, 2020.
- [29] M. Kryven, T. Ullman, B. Cowan, and J. Tenenbaum, “Plans or outcomes: How do we attribute intelligence to others?” 2021.
- [30] A. Netanyahu, T. Shu, B. Katz, A. Barbu, and J. B. Tenenbaum, “Phase: Physically-grounded abstract social events for machine social perception,” in *The AAAI Conference on Artificial Intelligence*, 2020.
- [31] C. Baker, R. Saxe, and J. Tenenbaum, “Bayesian theory of mind: Modeling joint belief-desire attribution,” in *Proceedings of the annual meeting of the cognitive science society*, vol. 33, no. 33, 2011.